

Computer Vision

Lectured by Wenjia Bai

Typed by Aris Zhu Yi Qing

March 21, 2022

Contents

1 Image Filtering	2	4.3 Laplacian of Gaussian (LOG)	8
1.1 Definition	2	4.4 Difference of Gaussian (DoG)	8
1.2 Common Filters	2	5 Feature Description	8
1.2.1 Moving Average (MA) Filter	2	5.1 Simple Descriptors	8
1.2.2 Identity Filter	3	5.2 Scale-Invariant Feature Transform (SIFT)	9
1.2.3 Gaussian Filter	3	5.3 Keypoint matching	9
1.2.4 Median Filter	3	5.4 RANSAC	10
1.3 Impulse Response	3	5.5 Acceleration	10
1.4 Convolution	4	5.5.1 Speeded-Up Robust Features (SURF)	10
2 Edge Detection	4	5.5.2 Binary Robust Independent Elementary Features (BRIEF)	10
2.1 Detection	4	5.5.3 Histograms of Oriented Gradient(HOG)	10
2.2 Edge Detection Filters	5	6 Image Classification	11
2.2.1 Prewitt Filter	5	6.1 Classification concepts	11
2.2.2 Sobel Filter	5	6.2 Support Vector Machine (SVN)	11
2.2.3 Magnitude and Orientation Calculation	5	6.3 Convolutional Neural Network	11
2.3 Canny Edge Detection	5	7 Object Detection	12
2.3.1 Criteria for Good Edge Detector	5	7.1 Approaches	12
2.3.2 Algorithm	5	7.2 Two-stage Object Detection Methods	12
3 Hough Transform	6	7.2.1 Selective Search	12
4 Interest Point Detection	7	7.2.2 Faster R-CNN	12
4.1 Definition	7	7.3 One-stage Object Detection Methods	13
4.2 Harris Detection	7		

8 Image Segmentation	13
8.1 Thresholding	13
8.2 K -means	13
8.3 Gaussian Mixture Model (GMM)	14
8.4 CNN	14
9 Motion	15
9.1 Definition	15
9.2 Optic flow constraint	15
9.3 Lucas-Kanade method	15
9.4 Horn-Schunck Method	16
9.5 Learning-based Flow Method	16
9.6 Correlation Filter	17
9.7 Tracking by Detection	17
9.8 Action Recognition	17

1 Image Filtering

1.1 Definition

- **Kernel**: a small matrix used to apply effects, e.g. blurring.
- **Separable kernel**: kernels that can be separated as two or more simple filters.
- **Padding**: The action of adding pixels around the borders (e.g. with value 0) so that applying filters will not reduce the size of the image.
- **Low-pass (smoothing) filter**: filters that keep the low-frequency signals, e.g. MA filter
- **High-pass (sharpening) filter**: filters that highlight the high-frequency signals, e.g. (identity + (identity - MA)) filter, or

$$\begin{pmatrix} -\frac{1}{8} & -\frac{1}{8} & -\frac{1}{8} \\ -\frac{1}{8} & 2 & -\frac{1}{8} \\ -\frac{1}{8} & -\frac{1}{8} & -\frac{1}{8} \end{pmatrix}.$$

- **Denoising filter**: filters to remove noise, e.g. median filter, non-local means, block-matching and 3D filtering (BM3D), etc.

1.2 Common Filters

1.2.1 Moving Average (MA) Filter

- In a 2D case, the MA kernel is a $\mathbb{R}^{K \times K}$ matrix in the following form

$$\frac{1}{K^2} \begin{pmatrix} 1 & 1 & \dots & 1 \\ 1 & 1 & \dots & 1 \\ \vdots & \vdots & \ddots & \vdots \\ 1 & 1 & \dots & 1 \end{pmatrix}$$

with a time complexity of $O(N^2 K^2)$, where N is the length of image.

- MA kernel is separable, for instance

$$\begin{pmatrix} \frac{1}{4} & \frac{1}{4} \\ \frac{1}{4} & \frac{1}{4} \end{pmatrix} = \begin{pmatrix} \frac{1}{2} & 1 \\ 1 & \frac{1}{2} \end{pmatrix} * \begin{pmatrix} \frac{1}{2} \\ \frac{1}{2} \end{pmatrix},$$

reducing the time complexity to $O(N^2K)$.

- Purpose:
 - remove high-frequency signal (noise or sharpness)
 - result in a smooth but blurry image

1.2.2 Identity Filter

The identity filter kernel is

$$\begin{pmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{pmatrix}$$

1.2.3 Gaussian Filter

- The Gaussian kernel is a 2D Gaussian distribution

$$h(i, j) = \frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}}$$

with $i, j = 0$ as the centre of the kernel.

- While its support is infinite, small values outside $[-k\sigma, k\sigma]$ can be ignored, e.g. $k = 3$ or $k = 4$.
- 2D Gaussian filter is separable with

$$h(i, j) = h_x(i) * h_y(j)$$

where

$$h_x(i) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{i^2}{2\sigma^2}},$$

because

$$\begin{aligned} f[x, y] * h[x, y] &= \sum_i \sum_j f[x - i, y - j] h[i, j] \\ &= \sum_i \sum_j f[x - i, y - j] \left(\frac{1}{2\pi\sigma^2} e^{-\frac{i^2+j^2}{2\sigma^2}} \right) \\ &= \sum_i \left(\sum_j f[x - i, y - j] \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{j^2}{2\sigma^2}} \right) \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{i^2}{2\sigma^2}} \end{aligned}$$

$$\begin{aligned} &= \sum_i (f * h_y)[x - i] \frac{1}{\sqrt{2\pi}} e^{-\frac{i^2}{2\sigma^2}} \\ &= (f * h_y) * h_x \end{aligned}$$

- Derivative of Gaussian filter h is

$$\frac{d(f * h)}{dx} = f * \frac{dh}{dx} = f * \frac{-x}{\sqrt{\pi}\sigma^3} e^{-\frac{x^2}{2\sigma^2}}.$$

Thus, the smaller the σ , the more detail in the magnitude map; larger σ suppresses noise and results in a smoother derivative. Different σ help find edges at different scale.

1.2.4 Median Filter

- non-linear
- often used for denoising
- Move the sliding window, and replace the centre pixel using the median value in the window.

1.3 Impulse Response

- For continuous signal, an impulse is a Dirac delta function $\delta(x)$, with

$$\delta(x) = \begin{cases} +\infty, & \text{if } x = 0 \\ 0, & \text{otherwise} \end{cases}$$

such that $\int_{-\infty}^{\infty} \delta(x) dx = 1$. For discrete signal, an impulse is a Kronecker delta function $\delta[i]$, with

$$\delta[i] = \begin{cases} 1, & \text{if } i = 0 \\ 0, & \text{otherwise.} \end{cases}$$

- The impulse response h is the output of a filter when the input is an impulse. It completely characterises a linear time-invariant filter.

- shifting the input signal k steps corresponds to the same output signal but shifted by k steps as well, e.g. assuming $f[n] = \delta[n]$, $g[n] = h[n]$,
 - * $g[n] = 10f[n]$ is time-invariant and amplifies the input by a constant.
 - * $g[n] = nf[n]$ is *not* time-invariant since the amount it amplifies the input depends on the
- if input $f_1[n]$ leads to $g_1[n]$, $f_2[n]$ leads to $g_2[n]$, we will have

$$\text{output}(\alpha f_1[n] + \beta f_2[n]) = \alpha g_1[n] + \beta g_2[n].$$

1.4 Convolution

- **Convolution:** output g can be described as the convolution between an input f and impulse response h as

$$g[n] = f[n] * h[n] = \begin{cases} \sum_{m=-\infty}^{\infty} f[m]h[n-m] & \text{discrete} \\ \int_{m=-\infty}^{\infty} f(m)h(n-m) & \text{continuous} \end{cases}$$

- Note that if we describe input signal $f[n]$ as

$$f[n] = \sum_{i=0}^n f[i]\delta[n-i]$$

and we know the output of $\delta[n]$ is $h[n]$, we can write the output as

$$g[n] = \sum_{i=0}^n f[i]h[n-i]$$

- commutative, i.e.

$$f[n] * h[n] = h[n] * f[n]$$

- associative, i.e.

$$f * (g * h) = (f * g) * h$$

- distributivity, i.e.

$$f * (g + h) = f * g + f * h \quad \text{and} \quad \frac{d(f * g)}{dx} = \frac{df}{dx} * g = f * \frac{dg}{dx}$$

- In 2D discrete case for image filtering,

$$\begin{aligned} g[m, n] &= f[m, n] * h[m, n] = \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f[i, j]h[m-i, n-j] \\ &= \sum_{i=-\infty}^{\infty} \sum_{j=-\infty}^{\infty} f[m-i, n-j]h[i, j] \end{aligned}$$

if the dimension of the kernel is $(2M+1) \times (2N+1)$, we can write

$$(f * h)[m, n] = \sum_{i=-M}^M \sum_{j=-N}^N f[m-i, n-j]h[i, j]$$

with $h[0, 0]$ being the centre of the filter, (m, n) being the location in the image which the kernel's center is on.

- If a big filter f_b can be separated into convolution g and h , we can first convolve with g , then h

$$f * f_b = f * (g * h) = (f * g) * h.$$

2 Edge Detection

2.1 Detection

	finite difference	convolution kernel
Forward difference	$f'[x] = f[x+1] - f[x]$	$[1, -1, 0]$
Backward difference	$f'[x] = f[x] - f[x-1]$	$[0, 1, -1]$
Central difference	$f'[x] = (f[x+1] - f[x-1])/2$	$[1, 0, -1]$

2.2 Edge Detection Filters

2.2.1 Prewitt Filter

Along the x -axis and the y -axis, we have respectively

$$\begin{pmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{pmatrix} \quad \text{and} \quad \begin{pmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{pmatrix}.$$

They are separable, i.e.

$$\begin{pmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix} * \begin{pmatrix} 1 & 0 & -1 \end{pmatrix}.$$

2.2.2 Sobel Filter

Along the x -axis and the y -axis, we have respectively

$$\begin{pmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{pmatrix} \quad \text{and} \quad \begin{pmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{pmatrix}.$$

They are also separable, i.e.

$$\begin{pmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{pmatrix} = \begin{pmatrix} 1 \\ 2 \\ 1 \end{pmatrix} * \begin{pmatrix} -1 & 0 & -1 \end{pmatrix}$$

2.2.3 Magnitude and Orientation Calculation

Let h_x denotes the horizontal filter, h_y denotes the vertical filter, we can compute the magnitude and the orientation as

$$\begin{aligned} g_x &= f * h_x && \text{derivative along } x\text{-axis} \\ g_y &= f * h_y && \text{derivative along } y\text{-axis} \\ g &= \sqrt{g_x^2 + g_y^2} && \text{magnitude of the gradient} \\ \theta &= \arctan(g_y, g_x) && \text{angle of the gradient} \end{aligned}$$

2.3 Canny Edge Detection

2.3.1 Criteria for Good Edge Detector

- good detection: low probability of FP/FN on marking edge points
- good localisation: mark as close as the centre of true edge
- single response: only one response to a single edge

2.3.2 Algorithm

1. perform Gaussian filtering to suppress noise
2. calculate the gradient magnitude $M(x, y)$ and direction
3. apply Non-Maximum Suppression (NMS) to get a single response for each edge

$$M(x, y) = \begin{cases} M(x, y) & \text{if local maximum} \\ 0 & \text{otherwise} \end{cases}$$

4. perform hysteresis thresholding to find potential edges with two thresholds t_{low} and t_{high}

$$\begin{cases} M(x, y) \geq t_{\text{high}} & \text{accept} \\ M(x, y) < t_{\text{low}} & \text{reject} \\ \text{Otherwise} & \text{iteratively check neighbouring pixels and} \\ & \text{accept if connected to an edge pixel.} \end{cases}$$

5. evaluation/performance

- good detection — FP reduced by Gaussian smoothing and FN reduced by hysteresis thresholding to find weak edges
- good localisation — NMS finds locations based on gradient magnitude and direction
- single response — NMS finds one single point in the neighbourhood




Figure 1: hough transform grid

3 Hough Transform

- **Hough transform** is a transform from image space to parameter space, e.g. from an edge map to the two parameters of a line.
- output is a parametric model, given the input edge points
- each edge point vote for possible models in the parameter space
- Example:
 - use slope intercept $b = y - mx$ to be the line model
 - each edge point poll vote for different b and m values (also lines in parameter space)
 - In practice, we use 2D bins to divide the parameter space; each point increases the vote by 1 in one of the bins, as shown in figure 1.
 - But the parameter space could be too large, $[-\infty, +\infty]$; use *normal form* instead

$$x \cos \theta + y \sin \theta = \rho$$

so at least for one dimension, $\theta \in [0, \pi)$.

- Line detection by Hough transform:
 1. Initialize the bins $H(\rho, \theta)$ to all zeros.
 2. For each (x, y) and each θ , calculate ρ and $H(\rho, \theta)++$.
 3. Find (ρ, θ) where $H(\rho, \theta)$ is a local maximum and larger than a threshold.
 - local maximum so there can be multiple solutions, to reduce FN
 - larger than a threshold so as to reduce FP
 4. The detected lines are given by $\rho = x \cos \theta + y \sin \theta$.
- robust to noise/occlusion
 - edge map is often generated after image smoothing
 - broken/unoccluded edge points can still vote and contribute to line detection
- Circle detection by Hough transform
 - parameter space also forms circles
 - If radius r is also unknown, then 3D parameter space $H(a, b, r)$
 - parameterize to be $x = a + r \cos \theta$ and $y = b + r \sin \theta$, since we know θ from edge detection, we can narrow the voting area to move along θ for a distance r .
- Pros and Cons:
 - + detects multiple instances
 - + robust to image noise
 - + robust to occlusion
 - Computational complexity is quite high. For each edge point, we need to vote to a 2D or even 3D parameter space.
 - need to carefully set parameters such as those for edge detectors, the threshold for the accumulator or range of circle radius.

4 Interest Point Detection

4.1 Definition

- **Interest point**: points that we are interested in and are useful for subsequent image processing and analysis. They are also called *key-points*, *landmarks*, *low-level* features.

4.2 Harris Detection

- small windows to tell if it is a corner
 - flat — change of intensity in neither direction
 - edge — change of intensity along just one direction
 - corner — change of intensity along both direction
- If the window shifts by $[u, v]$, together with Taylor expansion

$$I(x + u, y + v) = I(x, y) + uI_x(x, y) + vI_y(x, y) + \dots$$

the change of intensities is given by

$$\begin{aligned} E(u, v) &= \sum_{(x, y) \in W} w(x, y) [I(x + u, y + v) - I(x, y)]^2 \\ &= \sum w(x, y) [uI_x(x, y) + vI_y(x, y)]^2 \\ &= \sum w(x, y) (u^2 I_x^2 + 2uv I_x I_y + v^2 I_y^2) \\ &= \sum w(x, y) \begin{pmatrix} u & v \end{pmatrix} \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix} \\ &= \begin{pmatrix} u & v \end{pmatrix} \sum w(x, y) \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix} \\ &:= \begin{pmatrix} u & v \end{pmatrix} M \begin{pmatrix} u \\ v \end{pmatrix} \end{aligned}$$

- E is large if image derivatives are large
- We can infer the directions of change from M

- If $M = \begin{pmatrix} 0 & 0 \\ 0 & 0 \end{pmatrix}$, this is a flat region.
- If $M = \begin{pmatrix} 10 & 0 \\ 0 & 0.1 \end{pmatrix}$, large change if shift along $u \Rightarrow$ an edge.
- If $M = \begin{pmatrix} 10 & 0 \\ 0 & 10 \end{pmatrix}$, large change in whichever way \Rightarrow a corner.
- If more complicated M , we can do **eigen decomposition** to obtain diagonal matrix Λ , and compare against the above 3 matrices.

- **Cornerness** metrics:

- Harris and Stephens:

$$R = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2 = \det(M) - k(\text{trace}(M))^2$$

where k is a small number, $k = 0.05$.

- Kanade and Tomasi:

$$R = \min(\lambda_1, \lambda_2)$$

- Noble:

$$R = \frac{\lambda_1 \lambda_2}{\lambda_1 + \lambda_2 + \epsilon}$$

- also detects blobs and textures
- algorithm:

1. compute x and y derivative of an image

$$I_x = G_x * I, I_y = G_y * I$$

where G can be e.g. sobel filter

2. at each pixel, compute

$$M = \sum_{x, y} w(x, y) \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix}$$

3. calculate detector response and detect interest point which are local maxima and whose R is above threshold.

- scaled Harrison detector:

$$M = \sum_{x,y} w(x,y) \sigma^2 \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix}$$

so that

- at each scale σ , determine the scale which gives us the largest detector response
- prevent "the larger the scale, the smaller the derivative magnitude" problem
- Algorithm: perform original Harrison detector steps for each scale, and then determine the interest point.

4.3 Laplacian of Gaussian (LOG)

- First perform Gaussian smoothing, then Laplacian Operator, i.e.

$$\Delta(f * h) = \frac{\partial^2(f * h)}{\partial x^2} + \frac{\partial^2(f * h)}{\partial y^2} = f * \left(\frac{\partial^2 h}{\partial x^2} + \frac{\partial^2 h}{\partial y^2} \right)$$

where h is the Gaussian kernel and $\Delta f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$, with the following laplacian filter

$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix} = \begin{pmatrix} 0 & 0 & 0 \\ 1 & -2 & 1 \\ 0 & 0 & 0 \end{pmatrix} + \begin{pmatrix} 0 & 1 & 0 \\ 0 & -2 & 0 \\ 0 & 1 & 0 \end{pmatrix}$$

Since we know the formation of a 2D Gaussian, we can derive the following:

$$\frac{\partial^2 h}{\partial x^2} + \frac{\partial^2 h}{\partial y^2} = -\frac{1}{\pi \sigma^4} \left(1 - \frac{x^2 + y^2}{2\sigma^2} \right) e^{-\frac{x^2 + y^2}{2\sigma^2}}$$

- To ensure comparability between scales, we do

$$\text{LoG}_{\text{norm}}(x, y, \sigma) = \sigma^2 (I_{xx}(x, y, \sigma) + I_{yy}(x, y, \sigma))$$

similar to Harrison detector.

4.4 Difference of Gaussian (DoG)

- DoG is defined as

$$\text{DoG}(x, y, \sigma) := I * G(k\sigma) - I * G(\sigma) \approx (k-1)\sigma^2 \nabla^2 G(x, y, \sigma)$$

as an approximation of LoG_{norm} using Gaussian filters at different scales.

- DoG filters are used in SIFT, which is a pipeline for detecting and describing interest points.
- Also scale-invariant, and follow similar procedures as above.

5 Feature Description

5.1 Simple Descriptors

- Pixel intensity
 - sensitive to absolute intensity value — same object with different illumination can have different intensity
 - not very discriminative — a single pixel doesn't represent any local content
- Patch intensities
 - + represent local pattern
 - + perform well if the images are of similar intensities and roughly aligned
 - sensitive to absolute intensity value
 - not rotation-invariant
- Gradient orientation
 - + sensitive to intensity changes
 - + orientation is robust to scaling
 - not rotation-invariant
- Histogram

- + robust to rotation
- + robust to scaling
- sensitive to intensity changes

5.2 Scale-Invariant Feature Transform (SIFT)

- detects and describe local features in images
- transforms an image into a large set of interest points, each of which is described by a feature vector that is invariant to translation, scaling, and rotation.
- Algorithm:
 1. detection of scale-space extrema (minima/maxima)

search across scales and pixel locations, looking for interest points using DoG.
 2. keypoint localisation

refine the estimate coordinates by fitting a quadratic curve to the neighbouring pixels and compute the new extrema.

Denote the DoG response as $D(x, y, \sigma)$ or $D(\mathbf{x})$, by Taylor expansion we have

$$D(\mathbf{x} + \Delta\mathbf{x}) = D(\mathbf{x}) + \frac{\partial D}{\partial \mathbf{x}} \mathbf{x} + \frac{1}{2} \Delta\mathbf{x} \frac{\partial^2 D}{\partial \mathbf{x}^2} \Delta\mathbf{x}$$

$$\Rightarrow \frac{\partial D(\mathbf{x} + \Delta\mathbf{x})}{\partial \Delta\mathbf{x}} = \frac{\partial D}{\partial \mathbf{x}} + \frac{\partial^2 D}{\partial \mathbf{x}^2} \Delta\mathbf{x} = 0 \Rightarrow \Delta\mathbf{x} = - \left(\frac{\partial^2 D}{\partial \mathbf{x}^2} \right)^{-1} \frac{\partial D}{\partial \mathbf{x}}$$
 3. orientation assignment

determine the dominant orientation θ for a neighbourhood of a keypoint

 - (a) an orientation histogram with 36 bins covering 360 degrees is created
 - (b) each pixel votes for an orientation bin, weighted by the gradient magnitude
 - (c) keypoint will be assigned an orientation, which is the peak of the histogram

So now we have the location (x, y) , scale σ , and dominant orientation θ .

4. keypoint descriptor

Compute a histogram of gradient orientations. In practice, subregions are used and each subregion has 4x4 samples.

- each subregions has one orientation histogram
- these multiple histograms together describe what it looks like around the keypoint.
- each subregion has an orientation histogram with 8 bins in 8 directions.

If 16 subregions are used, the descriptor (feature vector) has a dimension of $128 = 16 \times 8$.

- robust to rotation, scaling, and changes in illumination
 - rotation — relative to dominant direction
 - scaling — draw samples from a window proportional to size
 - changes in illumination — using gradient orientations

5.3 Keypoint matching

Suppose we find a keypoint (x, y) in A that corresponds to a keypoint (u, v) in B . We assume that they are related with an *affine* transformation:

$$\begin{pmatrix} u \\ v \end{pmatrix} = \begin{pmatrix} m_1 & m_2 \\ m_3 & m_4 \end{pmatrix} \begin{pmatrix} x \\ y \end{pmatrix} + \begin{pmatrix} t_x \\ t_y \end{pmatrix}$$

with many pairs of corresponding keypoints, we can write the equation as:

$$\begin{pmatrix} x_1 & y_1 & 0 & 0 & 1 & 0 \\ 0 & 0 & x_1 & y_1 & 0 & 1 \\ x_2 & y_2 & 0 & 0 & 1 & 0 \\ 0 & 0 & x_2 & y_2 & 0 & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{pmatrix} \begin{pmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_x \\ t_y \end{pmatrix} = \begin{pmatrix} u_1 \\ v_1 \\ u_2 \\ v_2 \\ \vdots \end{pmatrix}$$

which can be written as a linear system

$$A\mathbf{m} = \mathbf{b}$$

where only \mathbf{m} is unknown, and solve the least-square problem as

$$\mathbf{m} = (A^T A)^{-1} A^T \mathbf{b}$$

5.4 RANSAC

However, the squared difference, $\|A\mathbf{m} - \mathbf{b}\|^2$, can be sensitive to outliers(noise) — points that are deemed to be corresponding but actually aren't. To ensure the robustness to outliers. RANSAC(Random Sample Consensus) deals with this.

1. randomly sample some points
2. fit a line along the sampled points
3. find the number of **inliers** within a threshold to the line
4. terminate if enough inliers have been found, or we have reached a certain number of iteration

5.5 Acceleration

5.5.1 Speeded-Up Robust Features (SURF)

- SURF only computes the gradients along horizontal and vertical directions using **Haar wavelets**.
- SURF applies very simple filters d_x and d_y

$$d_x = \begin{pmatrix} -1 & 1 \end{pmatrix}, d_y = \begin{pmatrix} -1^T \\ 1^T \end{pmatrix}$$

where $\mathbf{1}^T = (1 \ 1 \ \dots \ 1)$ called the Haar wavelets

- Summing pixel intensities with weights 1 and -1 is very fast, so the result for each subregion is defined by (the sum of values and the sum of absolute values)

$$\left(\sum d_x, \sum d_y, \sum |d_x|, \sum |d_y| \right)$$

- all four are low — homogeneous subregion
- zebra pattern/strips — $\sum |d_x|$ or $\sum |d_y|$ is high
- gradually increasing intensities — $\sum |d_x|$ and $\sum d_x$ are high
- Now we need $16 \times 4 = 64$ dimensions, and around 5 times faster than SIFT with Haar wavelets.

5.5.2 Binary Robust Independent Elementary Features (BRIEF)

- we compare an interest point p to another interest point q and get a binary value as output

$$\tau(p, q) = \begin{cases} 1 & \text{if } I(p) < I(q) \\ 0 & \text{otherwise.} \end{cases}$$

- randomly sample n_d pairs of points for binary tests; random pattern is only determined once and same pattern applied to all intersect points.
 - If $n_d = 256$, we have 256 tests in total, with a total of 32 bytes (SURF uses 64 bytes).
 - fast to compute with bit-shifting operator \ll and \gg .
 - use bitwise XOR between two descriptors to get the **hamming distance** instead of Euclidean distance, which is the bit count of the result.
- ignores rotation and scaling, assume images are taken from a moving camera with only translation.
- about 40-fold faster than SURF, which is around 200 times faster than SIFT.

5.5.3 Histograms of Oriented Gradient(HOG)

- we can extend the idea of feature descriptors to describe the feature of a large region, even a whole image.
- HOG divides a large region into a (dense) grid of cells, describes each cell, and concatenates the local descriptions to form a global description.
- The description vector \mathbf{v} for each cell is normalized to form a locally normalized descriptor as

$$\mathbf{v}_{\text{norm}} = \frac{\mathbf{v}}{\sqrt{\|\mathbf{v}\|_2^2 + \epsilon^2}}$$

with ϵ to ensure numerical stability

6 Image Classification

6.1 Classification concepts

See Intro to ML notes for k -NN, distance metrics, parameter tuning, cross-validation, stochastic gradient descent, multi-layer perceptron, forward/backward propagation

6.2 Support Vector Machine (SVN)

- If used as a *linear* SVN, the model is formulated as

$$\mathbf{w} \cdot \mathbf{x} + b = 0$$

The rule is to assign a class c to data x with

$$c = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b \geq 0 \\ -1 & \text{otherwise} \end{cases}$$

so we need to estimate \mathbf{w} and b to determine the decision boundary or *hyperplane*.

- The problem could be transformed to an optimisation problem of

$$\min_{\mathbf{w}, b} \|\mathbf{w}\|^2$$

subject to $y_i(\mathbf{w} \cdot \mathbf{x}_i + b) \geq 1$ for $i = 1, 2, \dots, N$. Or formulate as

$$\min_{\mathbf{w}, b} \|\mathbf{w}\|^2 + C \sum_{i=1}^N \max(0, 1 - y_i(\mathbf{w} \cdot \mathbf{x}_i + b))$$

since we cannot guarantee the data to be linearly separable. The second term of the above equation is called the **hinge loss**. The gradient of the loss function is

$$\nabla_{\mathbf{w}} L = 2\mathbf{w} - C \sum_{i=1}^N \nabla_{\mathbf{w}} h$$

where

$$\nabla_{\mathbf{w}} h = \begin{cases} -y_i \mathbf{x}_i & \text{if } y_i(\mathbf{w} \cdot \mathbf{x}_i + b) < 1 \\ 0 & \text{otherwise} \end{cases}$$

so at each iteration,

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} - \eta \nabla_{\mathbf{w}} L(\mathbf{w}^{(k)}, b^{(k)})$$

6.3 Convolutional Neural Network

- In CNN, each neuron only see a small local region in the layer before it, called the **receptive field**.
- e.g. A neuron can depend on a $5 \times 5 \times C$ cube of the input, where C refers to the number of color channels.
- More output neurons can be added as well, to form a $D \times 1 \times 1$ cube, where D denotes depth.
- Mathematically, a **convolution layer** is

$$a_d = f \left(\sum_{ijk} W_{dijk} x_{ijk} + b_d \right)$$

where d is the depth of output, ij is the dimension coordinate, and k is the input channel/depth.

- operations during convolution:
 - **padding**
 - **stride**
 - **dilation**: make a neuron look at a larger region (increase receptive field) by downsample image or increase kernel size
- **Pooling layer**:

- operation to make feature maps/representations smaller
- similar to image downsampling, e.g. if max pooling,

$$\begin{pmatrix} 1 & 2 \\ 5 & 6 \end{pmatrix} \Rightarrow (6)$$

pick the max element, and reduce the dimension

- issue:
 - *exploding* gradient — clip it, such as clip by value

$$\mathbf{g}_i = \begin{cases} v_{\min} & \text{if } g_i < v_{\min} \\ v_{\max} & \text{if } g_i > v_{\max} \\ \mathbf{g}_i & \text{otherwise} \end{cases}$$

or clip by norm

$$\mathbf{g} = \begin{cases} \frac{\mathbf{g}}{\|\mathbf{g}\|} v & \text{if } \|\mathbf{g}\| > v \\ \mathbf{g} & \text{otherwise} \end{cases}$$

- *vanishing gradient* — use different activation functions, e.g. ReLU, leaky ReLU, or
 - * Parametric ReLU

$$f(z) = \begin{cases} az & z < 0 \\ z & z \geq 0 \end{cases}$$

- * Exponential Linear Unit

$$f(z) = \begin{cases} a(e^z - 1) & z < 0 \\ z & z \geq 0 \end{cases}$$

- Representative CNN: LeNet, AlexNet, VGG

7 Object Detection

7.1 Approaches

- general ideas
 - at each sliding window, perform two tasks.
 - Firstly, classification (whether this is a cat or not)
 - secondly, localisation (bounding box coordinates) using CNN to predict bounding coordinates
 - however, too expensive to apply CNN on each pixel of the image
- two-stage detection:
- make an initial guess, propose some possible regions of interest
- classify what these regions are and predict the bounding boxes
- One-stage detection:
 - do not guess, but simply divide the image into grid cells
 - classify these grid cells and predict the bounding boxes

7.2 Two-stage Object Detection Methods

7.2.1 Selective Search

We can look at image features such as greyscale or gradients to separate the image into regions with similar features.

7.2.2 Faster R-CNN

RPN

- A Region Proposal Network (RPN) is used.
- The input image is fed into a CNN, gives the image's feature map.
- The feature map goes into a RPN, outputting interesting regions to be looked at.
- Using AlexNet/VGG-16 as backbone networks, using their last convolution layer as a feature map, we can use a 3×3 sliding window to perform binary classification at each location, giving it 0 if it isn't interesting, and 1 if it is.
- it handles objects of different sizes / aspect ratio by making k predictions (bounding boxes). e.g. 128^2 , 256^2 , 512^2 , $1:1$, $1:2$, $2:1$. These bounding boxes are called anchors.
- A bounding box can be described by its centre coordinates (x, y) and size (w, h) .
- For a convolution feature map of $W \times H$, $W \times H \times k$ anchors are predicted, and only the highest scoring boxes are kept.
- The loss is defined as

$$L(p, t) = \underbrace{\sum_{i=1}^{n_{\text{anchor}}} L_{\text{cls}}(p_i, p_i^*)}_{\text{classification loss}} + \lambda \underbrace{\sum_{i=1}^{n_{\text{anchor}}} 1_{y=1} L_{\text{loc}}(t_i, t_i^*)}_{\text{localisation loss}}$$

where $*$ denotes the ground truth

- we predict how to transform this anchor into the ground truth bounding box using:

$$\begin{aligned}
\text{anchor} &= (x_a, y_a, w_a, h_a) \\
\text{predicted bounding box} &= (x, y, w, h) \\
\text{predicted transformation} &= (t_x, t_y, t_w, t_h) \\
t_x &= \frac{x - x_a}{w_a} \\
t_y &= \frac{y - y_a}{h_a} \\
t_w &= \log \frac{w}{w_a} \\
t_h &= \log \frac{h}{h_a}
\end{aligned}$$

RoI

- Combining classification and localisation, we have features for each region known as **RoI** (Region of Interest). We can use an RoI pooling layer, which we can calculate RoI location and size on the feature map and convert to a fixed size to be provided to the classifier.
- For each RoI, the classifier predicts the label class and refines the bounding box estimate:

$$L(p, t) = L_{\text{cls}}(p, y) + \lambda \cdot 1_{y \geq 0} L_{\text{loc}}(t, t^*)$$

where y denotes the ground truth, and is now a multi-class classification problems.

Comparison between RPN and Detection Network

Please see table 1.

7.3 One-stage Object Detection Methods

- e.g. YOLO, SSD. The RPN is changed from using a binary value for the classification loss to using a multi-class classifier, which predicts the object for each anchor.
- Faster R-CNN is more accurate but slower, and vice versa for SSD, since RPN estimates the regions size before looking closely at features and then refining it.

Table 1: RPN v.s. RoI

RPN	Detection network
The input to RPN is a 3×3 window on conv5 feature map.	The input to detection network is a proposed region, thus contains more accurate features.
It is class-agnostic. It only checks whether this is an RoI or not.	It classifies the region into a number of classes.
RPN needs to use a lot of anchors, because we do not know what the object looks like yet	It does not need to use anchors. We already know a rough size from the proposal.

- Backbone networks (VGG/AlexNet) can be changed out to improve performance.

8 Image Segmentation

8.1 Thresholding

This converts a greyscale image to a binary label map. At each pixel, the label is defined as

$$f(x) = \begin{cases} 1 & \text{if } I(x) \geq \text{threshold} \\ 0 & \text{otherwise.} \end{cases}$$

This doesn't require any training data, and only needs threshold as the parameter.

8.2 K-means

- unsupervised method
- Represents each cluster by its centre. Each data point (pixel intensity) is associated to the nearest cluster centre.
- It can be formulated as

$$\min \sum_{k=1}^K \sum_{x \in C_k} (x - \mu_k)^2 \quad \text{or} \quad \min \sum_{k=1}^K \sum_x \delta_{x,k} (x - \mu_k)^2 \quad (1)$$

where $\delta_{x,k}$ denotes membership of x in cluster k , μ_k denotes the centre of cluster k .

- Algorithm
 1. initialize μ_k , where $k = 1, 2, \dots, K$.
 2. For each iteration
 - (a) compute $\delta_{x,k}$ for each data point, assigning x to the nearest cluster centre μ_k .
 - (b) update μ_k according to the membership $\delta_{x,k}$.
 - (c) repeat until $\delta_{x,k}$ no longer changes or the maximum number of iterations is reached.
 - To determine which K to use, plot equation (1) with different K , and use the “elbow” method to determine the best value.
 - clustering can also be performed base on
 - color similarity
 - position + color similarity
 - other features
- so that x becomes a feature vector instead of a scalar.

8.3 Gaussian Mixture Model (GMM)

- unsupervised method
- GMM performs a *soft* assignment by assuming a Gaussian distribution for each cluster, formulated as

$$P(y_j = k | x_j, \pi_k, \mu_k, \sigma_k) = \pi_k \cdot \frac{1}{\sqrt{2\pi}\sigma_k} e^{-\frac{(x_j - \mu_k)^2}{2\sigma_k^2}}$$


where x_j is the intensity of feature for data point j , π_k is the mixing component for class k , y_j denotes the group which point j belongs to.

- See intro to ML for detailed algorithm.

8.4 CNN


- supervised method
- we can e.g. remove the fully connected layers and directly infer pixel-wise classification results from the 13×13 feature map
- **Fully convolutional network**: all layers are convolutional layers.
- **Convolutionalization**: replaces the fully connected layers with convolutional layers
- Need to apply **upsampling** operation to obtain a pixel-wise prediction
 - e.g. transposed convolution. For $\mathbb{R}^{4 \times 4} \mapsto \mathbb{R}^{8 \times 8}$, we can have one input correspond to a 3×3 region in the output, with some weight, with a stride of 2.
 - convolution in 1D as matrix operation v.s. 2D, see figure 2.

Convolution in 1D



$$\begin{bmatrix} w_1 & w_2 & w_3 & 0 \\ 0 & w_1 & w_2 & w_3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \end{bmatrix}$$

Transposed convolution in 1D



$$\begin{bmatrix} w_1 & 0 \\ w_2 & w_1 \\ w_3 & w_2 \\ 0 & w_3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \\ y_4 \end{bmatrix}$$

The weight matrix is transposed.

Figure 2: convolution v.s. transposed convolution in Matrix operation

- Thus a classification network can be transformed to a segmentation network by replacing the fully connected layers with convolutional and transposed convolutional layers.
- **Mask R-CNN**

- the input is the convlution feature map
- fed into two branches: detection and segmentation, so performing two tasks at the same time
- segmentation branch is fully convolutional

9 Motion

9.1 Definition

- **Optic Flow:** The motion (flow) of brightness patterns (optic) in videos.
- The output of optic flow methods is a flow field, which describes the displacement vector for each pixel in the image.
- Assumption in optic flow:
 - brightness constancy
 - small motion
 - spatial coherence

9.2 Optic flow constraint

- Based on brightness constancy assumption, we have

$$I(x + u, y + v, t + 1) = I(x, y, t)$$

where I is the intensity, (x, y, t) is the spatial and temporal coordinate, and (u, v) is the displacement

- Based on small motion assumption, using first-order Taylor expansion we can approximate

$$\frac{\partial I}{\partial x}u + \frac{\partial I}{\partial y}v + \frac{\partial I}{\partial t} = 0 \quad \text{or} \quad I_x u + I_y v + I_t = 0$$

9.3 Lucas-Kanade method

- At each pixel p , we have the optic flow constraint equation

$$(I_x(p) \quad I_y(p)) \begin{pmatrix} u \\ v \end{pmatrix} = -I_t(p)$$

- For a small neighbourhood, e.g. a 3×3 window, we have a system of linear equations,

$$\begin{pmatrix} I_x(p_1) & I_y(p_1) \\ I_x(p_2) & I_y(p_2) \\ \vdots & \vdots \\ I_x(p_N) & I_y(p_N) \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix} = - \begin{pmatrix} I_t(p_1) \\ I_t(p_2) \\ \vdots \\ I_t(p_N) \end{pmatrix}$$

equivalently a problem of the form $Ax = b$.

- solution is given by

$$x = (A^T A)^{-1} A^T b \tag{2}$$

where

$$x = \begin{pmatrix} u \\ v \end{pmatrix}, A^T A = \sum_p \begin{pmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{pmatrix}, A^T b = - \sum_p \begin{pmatrix} I_x I_t \\ I_y I_t \end{pmatrix}$$

- Note that the matrix $A^T A$ also appears in the Harris detector 4.2.
- Using concepts in linear algebra, since

$$\text{cond}(A^T A) = \frac{\lambda_{\max}}{\lambda_{\min}}$$

for flat regions, we have large condition number \implies unstable calculation of the inverse; for corner regions, we have larger λ_{\min} , so smaller condition number \implies more stable inverse calculation.

Intuitively, the motion of flat/edge region couldn't be figured around a small neighbourhood accurately, but a corner is fine.

- **multi-scale:** motion at different scale. For each scale l from coarse to fine

- upsample the flow estimate at previous scale $l + 1$ with

$$u^{(l)} = 2u^{(l+1)}, v^{(l)} = 2v^{(l+1)}$$

- compute the warped image

$$J_{\text{warped}}^l = J^l(x + u^{(l)}, y + v^{(l)})$$

- for image I^l and J_{warped}^l at this scale, compute the image gradients I_x , I_y and I_t .
 - I_x and I_y are calculated from image I^l .
 - $I_t = J_{\text{warped}}^l(x, y) - I^l(x, y)$.
- estimate the incremental flow using

$$\begin{pmatrix} u^\delta \\ v^\delta \end{pmatrix} = (A^T A)^{-1} A^T b$$

- update the flow at scale l : $u^{(l)} = 2u^{(l+1)} + u^\delta$, $v^{(l)} = 2v^{(l+1)} + v^\delta$.

- For object tracking, the optimisation problem becomes

$$\min_{u,v} E(u, v) = \sum_x \sum_y [I(x, y) - J(x + u, y + v)]^2$$

where x, y are pixels in the **template image**, u, v indicates the motion from template I to image J . The solution is the same as equation (2).

9.4 Horn-Schunck Method

- an optimisation-based method, defining a global energy **functional** (a cost function) for the flow for all the pixels

$$\begin{aligned} E(u, v) &= \iint_{(x,y) \in \Omega} \underbrace{(I_x u + I_y v + I_t)^2}_{\text{optic flow constraint}} + \underbrace{\alpha}_{\text{weight}} \underbrace{(|\nabla u|^2 + |\nabla v|^2)}_{\text{smoothness term for the flow}} dx dy \\ &:= \iint_{(x,y) \in \Omega} L(x, y, u, v, u_x, u_y) dx dy \end{aligned}$$

where $u = u(x, y)$ and $v = v(x, y)$ are the unknown **functions** defined on the image pixels. They are solved by minimising the energy function $E(u, v)$.

- minimizing the functional comes down to solving its associated Euler-Lagrange equation,

$$\begin{aligned} \frac{\partial L}{\partial u} - \frac{\partial}{\partial x} \frac{\partial L}{\partial u_x} - \frac{\partial}{\partial y} \frac{\partial L}{\partial u_y} &= 0 \\ \frac{\partial L}{\partial v} - \frac{\partial}{\partial x} \frac{\partial L}{\partial v_x} - \frac{\partial}{\partial y} \frac{\partial L}{\partial v_y} &= 0 \end{aligned}$$

this follows that

$$\begin{aligned} (I_x u + I_y v + I_t) I_x - \alpha \Delta u &= 0 \\ (I_x u + I_y v + I_t) I_y - \alpha \Delta v &= 0 \end{aligned}$$

where $\Delta := \partial_{xx} + \partial_{yy}$.

- In practice, the Laplacian Δu can be approximated using finite differences

$$\Delta u = \bar{u} - u$$

where \bar{u} is the local average within a small neighbourhood.

- Plug in the approximation for Δu and reorganize, we obtain

$$\begin{aligned} (I_x^2 + \alpha)u + I_x I_y v &= \alpha \bar{u} - I_x I_t \\ I_x I_y u + (I_y^2 + \alpha)v &= \alpha \bar{v} - I_y I_t \end{aligned}$$

becoming an equation system for u and v , which can be solved as

$$\begin{aligned} u &= \bar{u} - \frac{I_x(I_x \bar{u} + I_y \bar{v} + I_t)}{I_x^2 + I_y^2 + \alpha} \\ v &= \bar{v} - \frac{I_y(I_x \bar{u} + I_y \bar{v} + I_t)}{I_x^2 + I_y^2 + \alpha}. \end{aligned}$$

thus u, v, \bar{u}, \bar{v} can be estimated iteratively.

9.5 Learning-based Flow Method

- A CNN takes in an input (two frames of an image), and outputs a flow field, visualized with a displacement vector at each pixel.
- The average End-point Error (EE) can be used to evaluate the performance of an optic flow method if the ground truth is available

$$EE = \frac{1}{N} \sqrt{(u(x, y) - u_{GT}(x, y))^2 + (v(x, y) - v_{GT}(x, y))^2}$$

9.6 Correlation Filter

- In a correlation filter, we aim to maximise the correlation between template features and image features in the next time frame

$$(f * h)[n] = \sum_{m=-\infty}^{\infty} f[m]h[n+m]$$

- The correlation can be more robust to illumination changes than sum of squared difference.
- In 2D images, we have a search window (in the next time frame) and a box with the same size as the template. We can find where the correlation achieves the maximum between the template features and features in a search window.
- The features can be learnt from CNNs.

9.7 Tracking by Detection

- tracking v.s. detection
 - tracking concerns about videos
 - In tracking, we have initial bounding box. We are only interested in how this object moves. We do not need to detect every object.
- online learning and offline learning
 - online: using information from this video, while it is streamed.
 - offline: using information from an existing training set.

Preferred method here is online learning.

- At each sliding window, we perform two tasks
 1. perform **binary** classification
 2. localisation to enable us to adjust the size of the object during tracking
- The positive and negative samples for binary classification are extracted from the first time frame. More samples can be generated by

- data augmentation: translation, rotation, etc.
- using more time frames while the video is streamed.

- Most layers of the convolutional network are pre-trained offline using large datasets such as ImageNet.
- Only the last few layers are trained online, using a lot of samples extracted from the video.

9.8 Action Recognition

We use both image and motion information as inputs

- image: static features extracted from a single image
- motion: dynamic features extracted from videos, e.g. optic flow fields

to train a classifier for action recognition.