

Distributed Systems

Lectured by Emil C Lupu

Typed by Aris Zhu Yi Qing

April 18, 2022

Contents

1 Characteristics

1.1	Distribution Transparencies	1
1.2	Challenges	1
1.3	Wrong Assumptions	2
1.4	Terminology	2

2 Architecture

2.1	Layered architecture	2
2.2	Object-based and service-oriented architectures	2
2.3	Message-based architectures	2
2.4	Peer-to-peer	2

3 Message-passing and IPC

1 Characteristics

1.1 Distribution Transparencies

1 Realize a coherent system by *hiding distribution* from the user where possible.

- **Access:** uniform access whether local or remote
- **Location:** access without knowledge of location
- **Concurrency:** sharing without interference (requires synchronization)
- **Replication:** hides use of redundancy (e.g. for fault tolerance)
- **Failure:** conceal failures by replication or recovery
- **Migration:** hides migration of components (e.g. for load balancing)
- **Performance:** hide performance variations (e.g. through use of scheduling and reconfiguration)
- **Scaling:** permits expansion by adding more resources (e.g. cloud)

1.2 Challenges

- **Heterogeneity:** different OS, data representation, implementations, etc.
- **Openness:** need to define *interfaces* for components to easily scale up systems
- **Security:** control access to preserve integrity and confidentiality
- **Concurrency:** inconsistencies may arise with interleaving requests
- **Failure handling:** transient/permanent failures could occur at any time. It is difficult detect them and to maintain consistency.
- **Scalability:** size of the system makes it difficult to maintain information about *system state*.

1.3 Wrong Assumptions

- The network is reliable, secure & homogeneous.
- The topology does not change.
- The latency is zero.
- The bandwidth is infinite.
- Transport cost is zero.
- There is one administrator.

1.4 Terminology

- **Client:** an entity initiating an interaction
- **Server:** a component responds to interactions usually implemented as a process
- **Service:** a component of a computer system that manages a collection of resources and presents their functionality to users.
- **Middleware:** software layer between the application and the OS masking the heterogeneity of the underlying system.

2 Architecture

2.1 Layered architecture

- e.g. Network stack. Control flows downwards, results flow upwards.
- + framework is simple and easy to learn and implement
- + reduced dependency due to layer separation
- + testing is easier with such modularity
- + cost overheads are fairly low
- scalability is difficult due to fixed framework structure
- difficult to maintain, since a change in a single layer can affect the entire system because it operates as a single unit
- parallel processing is not possible

2.2 Object-based and service-oriented architectures

- e.g. RMI
- + reusability, easy maintainability and greater reliability due to modularity
- + improved scalability and availability: multiple instances of a single service can run on different servers at the same time.
- Increased overhead: Service interactions require validations of inputs, thereby increasing the response time and machine load, and reducing the overall performance

2.3 Message-based architectures

	Temporally coupled	Temporally decoupled
Referentially coupled	Direct process messaging	Messaging via mailbox
Referentially decoupled	Event-based (publish-subscribe)	Shared data spaces

- Referentially coupled: processes name sender/receiver in their communication.
- Temporally coupled: both sender and receiver need to be up and running.

2.4 Peer-to-peer

- structured: Each node is indexed so that the location is known, and messages are routed according to the topology.
- unstructured: *flooding* or *random walks* or both.
- + no server needed since individual workstations are used to access files
- + resilient to computer failures, since it does not disrupt any other part of the network
- + very scalable
- poor performance with larger networks since each computer is being accessed by other users
- no central file system, hard to look up or backup
- ensuring that viruses are not introduced into the network is the responsibility of each individual user.
- There is no security other than assigning permissions.

3 Message-passing and IPC

- Asynchronous send: sender continues its execution once the message has been copied out of its address space
 - + mostly used with *blocked receive*
 - + underlying system must provide buffering for receiving messages independently of receiver processes
 - + *loose* coupling: sender does not know when message will be received, does not suspend execution until the message has been received
 - *Buffer exhaustion* (no flow control)
 - formal verification is more difficult, as need to account for the state of the buffers
- Synchronous send: *blocked send*, where the sender is held up until actual receipt of the message by the destination.
 - + usually used with blocking receive, where receiver execution is suspended until a message is received.
 - + synchronization between sender and receiver
 - + generally easier to formally reason about synchronous systems
 - what if no receivers? message loss?
 - No multi-destination, requiring synchronization with all receivers.
 - implementation more complicated
 - The underlying communication service is expected to be *reliable*, i.e. to guarantee in order message delivery.
- Asynchronous receive: process continues execution if there are no messages. hardly provided as primitives
- Blocked receive: the destination process blocks if no message is available, and receives it into a target variable when available.
- Please check the coursework for how UDP client/server is implemented in Java, e.g. how datagram, socket, port are used.